

linearreg

September 26, 2024

```
[15]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[19]: csv_file = 'AirQuality.csv'
df = pd.read_csv(csv_file, sep=',')
```

```
[6]: df = df.drop(columns=['Unnamed: 15', 'Unnamed: 16'], axis=1)
```

```
[23]: df = df.dropna()
```

```
[24]: dt_series = pd.Series(data = [item.split("/") [2] + "-" + item.split("/") [1] +
    ↪ "-" + item.split("/") [0] for item in df['Date']], index=df.index) + ' ' + pd.
    ↪ Series(data=[str(item).replace(".", ":") for item in df['Time']], index=df.
    ↪ index)
dt_series = pd.to_datetime(dt_series)
```

```
[25]: df = df.drop(columns=['Date', 'Time'], axis=1)
df.insert(loc=0, column='DateTime', value=dt_series)
```

```
[26]: year_series = dt_series.dt.year
```

```
[27]: month_series = dt_series.dt.month

day_series = dt_series.dt.day
```

```
[28]: day_name_series = dt_series.dt.day_name()
```

```
[29]: df['Year'] = year_series
df['Month'] = month_series
df['Day'] = day_series
df['Day Name'] = day_name_series
```

```
[30]: df = df.sort_values(by='DateTime')
```

```
[32]: def comma_to_period(series):
        new_series = pd.Series(data=[float(str(item).replace(',', '.')) for item in
        ↪series], index=df.index)
        return new_series

[34]: cols_to_correct = ['CO(GT)', 'C6H6(GT)', 'T', 'RH', 'AH']
        for col in cols_to_correct:
            df[col] = comma_to_period(df[col])

[35]: df = df.drop(columns=['NMHC(GT)', 'CO(GT)', 'NOx(GT)', 'NO2(GT)'], axis=1)

[36]: aq_2004_df = df[df['Year'] == 2004]
        aq_2005_df = df[df['Year'] == 2005]

[38]: for col in aq_2004_df.columns[1:-4]:
        median = aq_2004_df.loc[aq_2004_df[col] != -200, col].median()
        aq_2004_df[col] = aq_2004_df[col].replace(to_replace=-200, value=median)

[39]: for col in aq_2005_df.columns[1:-4]:
        median = aq_2005_df.loc[aq_2005_df[col] != -200, col].median()
        aq_2005_df[col] = aq_2005_df[col].replace(to_replace=-200, value=median)

[40]: group_2004_month = aq_2004_df.groupby(by='Month')
        group_2005_month = aq_2005_df.groupby(by='Month')

[41]: df = pd.concat([aq_2004_df, aq_2005_df])

[42]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9357 entries, 0 to 9356
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DateTime              9357 non-null  datetime64[ns]
1   Unnamed: 0            9357 non-null  int64
2   PT08.S1(CO)           9357 non-null  float64
3   C6H6(GT)              9357 non-null  float64
4   PT08.S2(NMHC)         9357 non-null  float64
5   PT08.S3(NOx)          9357 non-null  float64
6   PT08.S4(NO2)          9357 non-null  float64
7   PT08.S5(O3)           9357 non-null  float64
8   T                     9357 non-null  float64
9   RH                    9357 non-null  float64
10  AH                    9357 non-null  float64
11  Year                  9357 non-null  int32
12  Month                 9357 non-null  int32
```

```


13 Day          9357 non-null    int32
14 Day Name     9357 non-null    object
dtypes: datetime64[ns](1), float64(9), int32(3), int64(1), object(1)
memory usage: 1.0+ MB

```

```

[43]: corr_df = df.iloc[:, 1:-4].corr()
plt.figure(figsize = (10, 6), dpi = 96)
sns.heatmap(data = corr_df, annot = True) # 'annot=True' fills the R values in
↳the heatmap cells.
plt.show()

```



```

[44]: from sklearn.model_selection import train_test_split

X = df['T'] # Pandas DataFrame containing only feature variables
y = df['RH'] # Pandas Series containing the target variable

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.33,
↳random_state = 42)

```

```

[45]: def errors_product():
    prod = (X_train - X_train.mean()) * (y_train - y_train.mean())
    return prod

```


```
[46]: def squared_errors():
    sq_errors = (X_train - X_train.mean()) ** 2
    return sq_errors
```

```
[47]: slope = errors_product().sum() / squared_errors().sum()
    intercept = y_train.mean() - slope * X_train.mean()

    print(f"Slope: {slope} \nIntercept: {intercept}")
```

Slope: -1.112053910794772
Intercept: 69.69110324644876

```
[48]: plt.figure(figsize = (12, 5), dpi = 96)
    plt.title("Regression Line", fontsize = 16)
    plt.scatter(df['T'], df['RH'])
    plt.plot(df['T'], slope * df['T'] + intercept, color = 'r', linewidth = 2,
             label = '$y = -1.1120x + 69.6911$')
    plt.xlabel("Temperature")
    plt.ylabel("Relative Humidity")
    plt.legend()
    plt.show()
```



```
[49]: from sklearn.linear_model import LinearRegression

    X_train_reshaped = X_train.values.reshape(-1, 1)
    y_train_reshaped = y_train.values.reshape(-1, 1)
    X_test_reshaped = X_test.values.reshape(-1, 1)
    y_test_reshaped = y_test.values.reshape(-1, 1)
```

```

lin_reg = LinearRegression()
lin_reg.fit(X_train_reshaped, y_train_reshaped)

print("Coefficient of  $x$  (or slope) ==>", lin_reg.coef_)
print("Intercept ==>", lin_reg.intercept_)

```

```

Coefficient of  $x$  (or slope) ==> [[-1.11205391]]
Intercept ==> [69.69110325]

```

```

[50]: from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

y_train_pred = lin_reg.predict(X_train_reshaped)
y_test_pred = lin_reg.predict(X_test_reshaped)

print(f"Train Set\n{'-' * 50}")
print(f"R-squared: {r2_score(y_train_reshaped, y_train_pred):.3f}")
print(f"Mean Squared Error: {mean_squared_error(y_train_reshaped, y_train_pred):.3f}")
print(f"Root Mean Squared Error: {np.sqrt(mean_squared_error(y_train_reshaped, y_train_pred)):.3f}")
print(f"Mean Absolute Error: {mean_absolute_error(y_train_reshaped, y_train_pred):.3f}")

print(f"\n\nTest Set\n{'-' * 50}")
print(f"R-squared: {r2_score(y_test_reshaped, y_test_pred):.3f}")
print(f"Mean Squared Error: {mean_squared_error(y_test_reshaped, y_test_pred):.3f}")
print(f"Root Mean Squared Error: {np.sqrt(mean_squared_error(y_test_reshaped, y_test_pred)):.3f}")
print(f"Mean Absolute Error: {mean_absolute_error(y_test_reshaped, y_test_pred):.3f}")

```

Train Set

```

-----
R-squared: 0.325
Mean Squared Error: 195.281
Root Mean Squared Error: 13.974
Mean Absolute Error: 11.289

```

Test Set

```

-----
R-squared: 0.346
Mean Squared Error: 187.026
Root Mean Squared Error: 13.676
Mean Absolute Error: 11.150

```