

# untitled

October 3, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv("Heart.csv")
```

```
[3]: print("Number of records in each label are")
print(df['target'].value_counts())
```

```
Number of records in each label are
target
1      165
0      138
Name: count, dtype: int64
```

```
[4]: print("\nPercentage of records in each label are")
print(df['target'].value_counts() * 100 / df.shape[0], "\n")
```

```
Percentage of records in each label are
target
1      54.455446
0      45.544554
Name: count, dtype: float64
```

```
[5]: df.head()
```

```
[5]:   Unnamed: 0  age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  \
0           0   63   1   3     145    233   1         0     150     0
1           1   37   1   2     130    250   0         1     187     0
2           2   41   0   1     130    204   0         0     172     0
3           3   56   1   1     120    236   0         1     178     0
4           4   57   0   0     120    354   0         1     163     1

      oldpeak  slope  ca  thal  target
0         2.3     0   0     1       1
```


1	3.5	0	0	2	1
2	1.4	2	0	2	1
3	0.8	2	0	2	1
4	0.6	2	0	2	1

```
[6]: def sigmoid(x):
      return pd.Series(1 / ( 1 + np.exp(-x)))
```

```
[7]: df['chol'].describe()
```

```
[7]: count    303.000000
      mean     246.264026
      std       51.830751
      min     126.000000
      25%     211.000000
      50%     240.000000
      75%     274.500000
      max     564.000000
      Name: chol, dtype: float64
```

```
[8]: plt.figure(figsize = (12,4), dpi = 96)
      plt.title("Histogram for cholesterol values")
      plt.hist(df['chol'], bins = 'sturges', edgecolor = 'black')
      plt.show()
```



```
[9]: def standard_scalar(series):
      new_series = (series - series.mean()) / series.std()
      return new_series
      scaled_chol = standard_scalar(df['chol'])
```

```
[10]: plt.figure(figsize = (12,4))
       plt.title("Histogram for cholesterol values")
```

```
plt.hist(scaled_chol, bins = 'sturges', edgecolor = 'black')
plt.show()
```



```
[11]: chol_sig_output = sigmoid(df['chol'])
chol_sig_output.describe()
```

```
[11]: count    303.0
      mean      1.0
      std       0.0
      min      1.0
      25%      1.0
      50%      1.0
      75%      1.0
      max      1.0
      Name: chol, dtype: float64
```


```
[12]: scaled_chol_sig_output = sigmoid(scaled_chol)
scaled_chol_sig_output.describe()
```

```
[12]: count    303.000000
      mean     0.492837
      std     0.198175
      min     0.089454
      25%     0.336179
      50%     0.469823
      75%     0.632919
      max     0.997829
      Name: chol, dtype: float64
```

```
[13]: def predict(sig_output, threshold):
      y_pred = [ 1 if output >= threshold else 0 for output in sig_output]
      return pd.Series(y_pred)
```

```
[14]: threshold = 0.5
heart_disease_pred = predict(scaled_chol_sig_output, threshold)
```

```
[15]: plt.figure(figsize=(13,3), dpi = 96)
plt.scatter(scaled_chol, heart_disease_pred)
plt.axhline(y = threshold, label = f'y = { threshold }', color = 'r')
plt.legend()
plt.show()
```



```
[16]: print(f"Threshold value: {threshold}")
print(f"\nPredicted value counts:\n{heart_disease_pred.value_counts()}")
print(f"\nActual value counts:\n{df['target'].value_counts()}")
```

Threshold value: 0.5

Predicted value counts:

0 167

1 136

Name: count, dtype: int64

Actual value counts:

target

1 165

0 138

Name: count, dtype: int64

```
[17]: from sklearn.metrics import confusion_matrix
```

```
[18]: print(confusion_matrix(df['target'], heart_disease_pred))
```

```
[[ 65  73]
```

```
 [102  63]]
```

```
[19]: from sklearn.metrics import classification_report
```

```
[20]: print(classification_report(df['target'], heart_disease_pred))
```

	precision	recall	f1-score	support
0	0.39	0.47	0.43	138
1	0.46	0.38	0.42	165
accuracy			0.42	303
macro avg	0.43	0.43	0.42	303
weighted avg	0.43	0.42	0.42	303

```
[21]: from sklearn.model_selection import train_test_split
```

```
[22]: X = df.drop(columns = 'target')
      y = df['target']
```

```
[23]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
      ↪random_state = 42)
```

```
[24]: from sklearn.linear_model import LogisticRegression
```

```
[25]: log_clf_1 = LogisticRegression()
      log_clf_1.fit(X_train, y_train)
      print(log_clf_1.score(X_train, y_train))
```

1.0

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[26]: y_train_pred = log_clf_1.predict(X_train)
```

```
[27]: print("\n Confusion Matrix \n")
      print(confusion_matrix(y_train, y_train_pred))
```

Confusion Matrix

```
[[ 97   0]
 [  0 115]]
```

```
[28]: print("\n Classification Report\n")
      print(classification_report(y_train, y_train_pred))
```

Classification Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	97
1	1.00	1.00	1.00	115
accuracy			1.00	212
macro avg	1.00	1.00	1.00	212
weighted avg	1.00	1.00	1.00	212

```
[29]: y_test_pred = log_clf_1.predict(X_test)
```

```
[30]: print(f"{'Test Set'.upper()}\n{'-' * 75}\nConfusion Matrix:")
      print(confusion_matrix(y_test, y_test_pred))
```

TEST SET

-----

Confusion Matrix:

```
[[40  1]
 [ 1 49]]
```

```
[31]: print("\nClassification Report")
      print(classification_report(y_test, y_test_pred))
```

Classification Report

	precision	recall	f1-score	support
0	0.98	0.98	0.98	41
1	0.98	0.98	0.98	50
accuracy			0.98	91
macro avg	0.98	0.98	0.98	91
weighted avg	0.98	0.98	0.98	91

```
[32]: def standard_scaler(series):
      new_series = (series - series.mean()) / series.std()
      return new_series
```

```
[33]: norm_X_train = X_train.apply(standard_scaler, axis = 0)
norm_X_test = X_test.apply(standard_scaler, axis = 0)
```

```
[34]: norm_X_train.describe()
norm_X_test.describe()
```

```
[34]:
```

	Unnamed: 0	age	sex	cp	trestbps	\
count	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01	
mean	-1.249001e-16	-2.488852e-16	-1.146824e-16	-9.760202e-18	-7.051746e-16	
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	
min	-1.644808e+00	-2.301763e+00	-1.661622e+00	-8.425578e-01	-1.853721e+00	
25%	-8.333169e-01	-8.354271e-01	-1.661622e+00	-8.425578e-01	-6.650121e-01	
50%	-6.722410e-02	1.797284e-01	5.952080e-01	-8.425578e-01	-1.662530e-02	
75%	9.258592e-01	6.309086e-01	5.952080e-01	1.123410e+00	4.696648e-01	
max	1.703301e+00	2.435630e+00	5.952080e-01	2.106394e+00	3.549502e+00	

	chol	fbs	restecg	thalach	exang	\
count	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01	
mean	-4.148086e-17	-5.490114e-17	0.000000	-5.343711e-16	-4.880101e-18	
std	1.000000e+00	1.000000e+00	1.000000	1.000000e+00	1.000000e+00	
min	-2.624853e+00	-4.938276e-01	-0.943037	-3.319275e+00	-7.148350e-01	
25%	-7.201088e-01	-4.938276e-01	-0.943037	-6.418709e-01	-7.148350e-01	
50%	-1.836075e-02	-4.938276e-01	-0.943037	1.078023e-01	-7.148350e-01	
75%	6.165541e-01	-4.938276e-01	0.963994	6.432832e-01	1.383552e+00	
max	3.679740e+00	2.002745e+00	2.871025	1.864180e+00	1.383552e+00	

	oldpeak	slope	ca	thal
count	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01
mean	1.082808e-16	-9.516197e-17	-1.464030e-17	1.903239e-16
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-8.367971e-01	-2.184053e+00	-8.102615e-01	-3.491486e+00
25%	-8.367971e-01	-5.812398e-01	-8.102615e-01	-4.364358e-01
50%	-3.799059e-01	-5.812398e-01	-8.102615e-01	-4.364358e-01
75%	5.719508e-01	1.021573e+00	9.246514e-01	1.091089e+00
max	3.884412e+00	1.021573e+00	2.659564e+00	1.091089e+00

```
[35]: from sklearn.feature_selection import RFE
from sklearn.metrics import f1_score
from sklearn.linear_model import LogisticRegression
```

```
[36]: dict_rfe = {}
```

```
[37]: for i in range(1, len(X_train.columns) + 1):
lg_clf_2 = LogisticRegression()
rfe = RFE(lg_clf_2, n_features_to_select=i)
rfe.fit(norm_X_train, y_train)
```

```

rfe_features = list(norm_X_train.columns[rfe.support_])
rfe_X_train = norm_X_train[rfe_features]

lg_clf_3 = LogisticRegression()
lg_clf_3.fit(rfe_X_train, y_train)

y_test_pred = lg_clf_3.predict(norm_X_test[rfe_features])

f1_scores_array = f1_score(y_test, y_test_pred, average = None)
dict_rfe[i] = {"features": list(rfe_features), "f1_score": f1_scores_array}
dict_rfe

```

```

[37]: {1: {'features': ['Unnamed: 0'], 'f1_score': array([0.98795181, 0.98989899])},
      2: {'features': ['Unnamed: 0', 'oldpeak'],
          'f1_score': array([0.98765432, 0.99009901])},
      3: {'features': ['Unnamed: 0', 'exang', 'oldpeak'],
          'f1_score': array([0.98765432, 0.99009901])},
      4: {'features': ['Unnamed: 0', 'exang', 'oldpeak', 'thal'],
          'f1_score': array([0.97560976, 0.98      ])},
      5: {'features': ['Unnamed: 0', 'restecg', 'exang', 'oldpeak', 'thal'],
          'f1_score': array([0.97560976, 0.98      ])},
      6: {'features': ['Unnamed: 0', 'sex', 'restecg', 'exang', 'oldpeak', 'thal'],
          'f1_score': array([0.96385542, 0.96969697])},
      7: {'features': ['Unnamed: 0',
                      'sex',
                      'cp',
                      'restecg',
                      'exang',
                      'oldpeak',
                      'thal'],
          'f1_score': array([0.97560976, 0.98      ])},
      8: {'features': ['Unnamed: 0',
                      'sex',
                      'cp',
                      'restecg',
                      'exang',
                      'oldpeak',
                      'ca',
                      'thal'],
          'f1_score': array([0.96385542, 0.96969697])},
      9: {'features': ['Unnamed: 0',
                      'sex',
                      'cp',
                      'restecg',
                      'exang',
                      'oldpeak',
                      'slope'],

```



```

    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
10: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'chol',
    'restecg',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
11: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'trestbps',
    'chol',
    'restecg',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
12: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'trestbps',
    'chol',
    'fbs',
    'restecg',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
13: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'trestbps',
    'chol',
    'fbs',
    'restecg',
    'thalach',

```

```

    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
14: {'features': ['Unnamed: 0',
    'age',
    'sex',
    'cp',
    'trestbps',
    'chol',
    'fbs',
    'restecg',
    'thalach',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])}]

```

```

[38]: pd.options.display.max_colwidth = 100
f1_df = pd.DataFrame.from_dict(dict_rfe, orient = 'index')
f1_df

```

```

[38]:          features \
1
[Unnamed: 0]
2
[Unnamed: 0, oldpeak]
3
[Unnamed: 0, exang, oldpeak]
4
exang, oldpeak, thal]          [Unnamed: 0,
5
exang, oldpeak, thal]          [Unnamed: 0, restecg,
6
exang, oldpeak, thal]          [Unnamed: 0, sex, restecg,
7
exang, oldpeak, thal]          [Unnamed: 0, sex, cp, restecg,
8
exang, oldpeak, ca, thal]      [Unnamed: 0, sex, cp, restecg,
9
oldpeak, slope, ca, thal]      [Unnamed: 0, sex, cp, restecg, exang,
10
oldpeak, slope, ca, thal]      [Unnamed: 0, sex, cp, chol, restecg, exang,

```

```

11             [Unnamed: 0, sex, cp, trestbps, chol, restecg, exang,
oldpeak, slope, ca, thal]
12             [Unnamed: 0, sex, cp, trestbps, chol, fbs, restecg, exang,
oldpeak, slope, ca, thal]
13             [Unnamed: 0, sex, cp, trestbps, chol, fbs, restecg, thalach, exang,
oldpeak, slope, ca, thal]
14 [Unnamed: 0, age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang,
oldpeak, slope, ca, thal]

```

```

                                f1_score
1      [0.9879518072289156, 0.98989898989899]
2      [0.9876543209876543, 0.9900990099009901]
3      [0.9876543209876543, 0.9900990099009901]
4              [0.975609756097561, 0.98]
5              [0.975609756097561, 0.98]
6      [0.963855421686747, 0.9696969696969697]
7              [0.975609756097561, 0.98]
8      [0.963855421686747, 0.9696969696969697]
9      [0.963855421686747, 0.9696969696969697]
10     [0.963855421686747, 0.9696969696969697]
11     [0.963855421686747, 0.9696969696969697]
12     [0.963855421686747, 0.9696969696969697]
13     [0.963855421686747, 0.9696969696969697]
14     [0.963855421686747, 0.9696969696969697]

```

```

[39]: lg_clf_4 = LogisticRegression()
      rfe = RFE(lg_clf_4, n_features_to_select = 3)

```

```

[40]: rfe.fit(norm_X_train, y_train)
      rfe_features = norm_X_train.columns[rfe.support_]
      print(rfe_features)
      final_X_train = norm_X_train[rfe_features]

```

```

Index(['Unnamed: 0', 'exang', 'oldpeak'], dtype='object')

```

```

[ ]:

```